

Highlights

From Detection to Maintenance Action: Streaming Sensor-Fault Typing for Self-Supervised Anomaly Detection in Industrial Systems

Marek Wadinger, Michal Kvasnica

- Streaming detector attributes each flagged sensor to a maintenance-relevant fault type
- Reaches 90% real-fault detection at 7.4% healthy false-positive rate on battery-depletion faults
- $O(1)$ -memory conditional-residual typing composed with a self-supervised online detector
- Operating envelope characterized: bias absorption, dead-band, and correlation limits made explicit
- Online adaptation is mandatory; a static prefix model false-alarms at 0.97 on a diurnal stream

From Detection to Maintenance Action: Streaming Sensor-Fault Typing for Self-Supervised Anomaly Detection in Industrial Systems

Marek Wadinger^{a,*}, Michal Kvasnica^a

^a*Institute of Information Engineering, Automation and Mathematics, Slovak University of Technology in Bratislava, Radlinskeho 9, 812 37, Bratislava, Slovakia*

Abstract

Streaming anomaly detectors tell operators *that* a sensor reading is abnormal and, when interpretable, *which* signal and in *which* direction. They do not tell operators *what to do*: whether to recalibrate, replace, or clean a sensor, or to investigate the process itself. We close this gap with a streaming sensor-fault classifier that attributes each flagged signal to a maintenance-relevant fault type—bias, drift, accuracy loss, or freezing—directly from the conditional residuals of a self-supervised online detector, using $O(1)$ memory and no history buffers. On real battery-depletion faults in the Intel Berkeley Lab deployment, the composed system reaches 90% real-fault detection at a 7.4% healthy false-positive rate, and types the resulting faults predominantly as freezing (0.93), the signature of sensors stuck at degraded values. We further characterize the operating envelope on controlled injections: freezing and drift are recovered with recall 1.000 and 0.741, while bias recall is monotone in conditional standard deviation—the interpretable conditional model absorbs constant offsets on strongly coupled signals (recall 0.059) and isolates them on weakly coupled ones (0.941). We show that two original claims do not survive validation: freezing is not a detection blind spot of the base detector (it catches every injected freeze, usually faster than the dedicated test), and adaptation robustness comes from graded change-point suppression rather than residual cancellation. The result is a deployable typing layer with a

**Phone numbers:* +421 902 810 324 (Marek Wadinger)

Email addresses: `marek.wadinger@stuba.sk` (Marek Wadinger),
`micHAL.kvasnica@stuba.sk` (Michal Kvasnica)

URL: `uiam.sk/~wadinger` (Marek Wadinger)

rigorously mapped—rather than overstated—operating envelope.

Keywords: Fault diagnosis, Sensor fault classification, Online learning, Anomaly detection, Self-supervised learning, Streaming algorithms

1. Introduction

In industrial monitoring, the cost of a sensor fault is paid not when it is detected but when it is acted upon. A reading that drifts away from its true value, a transmitter stuck at its last sample, or an instrument whose noise floor has grown all demand different maintenance responses: recalibration, replacement, cleaning, or process investigation. An operator who learns only that a signal is anomalous, and even which signal and in which direction, must still diagnose the failure mode before dispatching the correct action. The fault *type* is what drives the maintenance action, and it is precisely what most streaming detectors leave to human triage.

The base detector we extend is a self-supervised, online anomaly detector for SCADA-based industrial systems (Wadinger and Kvasnica, 2023). It scores each signal by the conditional cumulative probability of its current value given the remaining signals, flags a value when that probability falls below α or above $1 - \alpha$ (with $\alpha \approx 0.00133$, i.e. roughly three conditional standard deviations), and adapts its conditional model online with change-point awareness. Crucially, it is interpretable: it isolates the responsible signal and the direction of the excursion, and exposes per-signal operating limits. What it does not provide is the fault *type*—it localizes the anomaly but does not characterize the failure mode, leaving the detection-to-action gap open.

We close that gap with a streaming sensor-fault classifier that attributes each flagged signal to one of four maintenance-relevant types—bias, drift, accuracy loss, or freezing—computed from the same conditional residuals the base detector already produces. The classifier maintains only a small set of exponentially weighted running statistics per signal: it uses $O(1)$ memory, keeps no history buffers, and adds negligible cost to the streaming pipeline. The fault taxonomy itself is not new; we adopt the bias / noise / constant / drift categories established for sensor faults in wireless and environmental sensing (Sharma et al., 2010; Ni et al., 2009). Our contribution is the *streaming conditional-residual realization* of this taxonomy, composed directly with a self-supervised detector so that type attribution rides on the same online statistics as detection.

We lead with the practical outcome. On real battery-depletion faults from the Intel Berkeley Lab deployment (Bodik et al., 2004)—where a mote’s failing battery degrades its temperature and humidity readings to physically impossible values—the composed system reaches 90% real-fault detection at a 7.4% healthy false-positive rate, and types the resulting faults predominantly as freezing (0.93), the expected signature of a transmitter stuck at a degraded value. Online adaptation is not optional here: a static model fit on a short prefix of the same diurnal stream false-alarms at a rate of 0.97. The typing layer thus converts an adaptive detector’s alarm into an actionable maintenance label on a real fault population.

We then characterize the operating envelope honestly, treating it as support for—not a qualification of—the practical result. Two claims we initially expected to hold do not survive our own validation, and we report them as refuted. First, freezing is *not* a detection blind spot of the base detector: it catches every injected freeze on every channel, usually faster than a dedicated fixed-window freeze test, so the freeze test contributes type *attribution* (a stuck sensor versus a process excursion), not detection capability. Second, the robustness of the adaptive detector to faults that arrive during adaptation comes from graded change-point suppression, not from conditional-residual cancellation. We treat this honesty as a feature: the envelope is mapped rather than overstated, and the limitations—bias absorption on strongly coupled signals, the limited cancellation of a two-signal conditional model, and the masking dead-band introduced by suppression—are quantified and explained.

The contributions of this paper are:

1. A streaming, $O(1)$ -memory sensor-fault classifier that types each flagged signal as normal, bias, drift, accuracy loss, or freezing from the conditional residuals of a self-supervised online detector, with no history buffers (Section 3).
2. A demonstration on real battery-depletion faults that the composed system reaches 90% real-fault detection at a 7.4% healthy false-positive rate and assigns maintenance-relevant types, with online adaptation shown to be mandatory (Section 5).
3. A rigorous characterization of the operating envelope on controlled injections, including the monotone dependence of bias recall on conditional coupling, and two refuted claims (freezing as a blind spot; residual cancellation as the source of adaptation robustness) reported as

70 such (Sections 5 and 6).

71 2. Related Work

72 2.1. The base detector

73 The detector we extend scores anomalies in SCADA-based industrial sys-
74 tems by self-supervised online learning (Wadinger and Kvasnica, 2023). For
75 each signal it evaluates the conditional cumulative distribution of the current
76 value given the remaining signals under a multivariate Gaussian model, flags
77 the value when this probability lies outside $[\alpha, 1 - \alpha]$ with $\alpha \approx 0.00133$, and
78 updates the conditional model online while remaining aware of change points
79 so that legitimate regime shifts are absorbed rather than alarmed. The de-
80 tector is interpretable: it identifies the responsible signal, the direction of the
81 excursion, and exposes per-signal operating limits. It localizes anomalies but
82 does not classify the failure mode, which is the gap this work addresses. We
83 reuse its conditional-residual stream as the substrate for typing rather than
84 introducing a separate diagnostic pipeline.

85 2.2. Sensor-fault taxonomies

86 The categories we attribute—bias, drift, accuracy loss (noise), and freez-
87 ing (constant)—are standard in the sensor-fault literature rather than novel
88 to this work. Sharma et al. (2010) characterize the recurrent fault signatures
89 of deployed sensors, including offset (bias), gain, and stuck-at faults, and
90 study detection methods against them. Ni et al. (2009) likewise organize
91 sensor data faults into a taxonomy spanning calibration offset, drift, and
92 constant (frozen) readings. We adopt this established taxonomy unchanged;
93 our contribution is its streaming, conditional-residual realization, not the
94 categories themselves.

95 2.3. Fault-injection protocols

96 Because no public stream simultaneously labels all four fault types per
97 signal and per timestep, controlled injection is the established route to per-
98 type evaluation. Lai et al. (2021) introduce a benchmark methodology that
99 injects parameterized anomaly patterns into time series to evaluate detectors
100 under known ground truth. We follow this precedent, injecting the four
101 typed faults into a multivariate process at controlled magnitudes and onsets
102 so that recall and confusion can be measured per type, and we triangulate the
103 synthetic findings against the real-fault deployment described in Section 4.

104 2.4. Offline fault classifiers

105 A large body of work classifies sensor and process faults offline, typically
106 by training supervised or windowed models on labelled fault episodes and
107 extracting features over batches of history. Such methods can exploit rich
108 temporal context but require stored windows, labelled fault data, and a train-
109 ing phase, which conflicts with the streaming, self-supervised, and label-free
110 setting of the base detector. Our classifier instead derives the type from the
111 running conditional residuals already maintained by the detector, in $O(1)$
112 memory and without history buffers, so that typing is available at the same
113 latency and operational footprint as detection.

114 3. Method

115 3.1. Conditional residuals as the typing substrate

116 The base detector (Wadinger and Kvasnica, 2023) maintains, for each
117 signal i , a conditional model of its current value x_i given the remaining
118 signals x_{-i} . We define the conditional residual

$$r_i = x_i - \mu_{\text{cond},i}(x_{-i}), \quad (1)$$

119 the difference between the observed value and its conditional mean, and the
120 conditional standard deviation $\sigma_{\text{cond},i}$ already used by the detector to form
121 its score. Both quantities are available at every step at no additional mod-
122 elling cost. For brevity we write σ for the per-channel conditional standard
123 deviation $\sigma_{\text{cond},i}$ in the results, and we also use σ as a *magnitude unit*—a
124 fault or shift “of $k\sigma$ ” means k multiples of the relevant conditional standard
125 deviation; the intended sense is clear from context. The classifier operates
126 entirely on the stream of residuals $\{r_i\}$ and their conditional scale $\sigma_{\text{cond},i}$: it
127 never stores past samples and never refits a model over a window. This keeps
128 the per-signal state to a fixed set of scalars and the per-step cost constant,
129 matching the streaming, $O(1)$ -memory footprint of the detector it extends.

130 3.2. Running statistics

131 For each signal the classifier maintains a small set of exponentially weighted
132 moving (EWM) statistics over the residual stream. We track a fast (short
133 half-life) and a slow (long half-life) EWM mean of r_i , denoted $m_{\text{fast},i}$ and
134 $m_{\text{slow},i}$, and an EWM variance v_i of r_i . We additionally track the EWM

135 variance of the raw signal increments to detect a stalled transmitter. All up-
 136 dates are the standard constant-memory recursions; no buffer of past values
 137 is required. The half-lives are the only temporal hyperparameters of the typ-
 138 ing layer and are shared across signals unless per-signal calibration is applied
 139 (Section 6).

140 3.3. The four type tests

141 Given the running statistics, a flagged signal is attributed to a type by
 142 the following tests, each expressed in units of the conditional scale so that
 143 thresholds are dimensionless.

144 *Bias..* A persistent offset of the residual from zero, with stable variance,
 145 indicates a constant calibration error. We declare *bias* when the slow residual
 146 mean is large in conditional-scale units, $|m_{\text{slow},i}|/\sigma_{\text{cond},i}$ exceeding a threshold,
 147 while the residual variance remains close to its baseline.

148 *Drift..* A gradually growing offset is distinguished from a static one by the
 149 gap between the short- and long-horizon EWM means. We declare *drift* when
 150 $|m_{\text{fast},i} - m_{\text{slow},i}|/\sigma_{\text{cond},i}$ exceeds a threshold, i.e. when the recent residual
 151 mean is moving away from its longer-run value. We name this honestly: it
 152 is a short-versus-long EWM mean-gap statistic, *not* a fitted regression slope.
 153 It requires no stored window and no least-squares fit; it responds to the
 154 direction and recency of residual movement rather than estimating a rate.

155 *Accuracy loss..* An inflated noise floor with an unchanged mean indicates
 156 degraded measurement accuracy. We declare *accuracy loss* when the residual
 157 variance v_i exceeds its baseline by a multiplicative factor while the slow
 158 residual mean stays near zero.

159 *Freezing..* A transmitter stuck at its last value produces near-zero signal in-
 160 crements while the conditional expectation continues to move. We declare
 161 *freezing* when the EWM variance of the raw increments collapses below an
 162 absolute floor `freeze_eps` and below a fraction `freeze_var_ratio` of the sig-
 163 nal’s marginal variance, scaled by `freeze_abs_scale`. These three constants
 164 govern how aggressively a quiet signal is judged frozen; their calibration is
 165 examined in Section 5 (the freeze-on-quiescent ablation). A signal that satis-
 166 fies none of the tests while flagged is labelled *normal* with respect to typing
 167 (e.g. a peer signal pulled by cross-talk), and an unflagged signal is *normal*
 168 by default.

169 3.4. Change-point graded suppression

170 A fault that arrives while the base detector is adapting to a genuine
171 regime shift can be either masked by the adaptation or mistaken for one.
172 The detector’s change-point machinery applies a graded suppression to the
173 typing decision around detected change points, controlled by a scale param-
174 eter `suppress_scale`: at `scale = 1` no suppression is applied, the default
175 `scale = 5` applies graded suppression, and `scale = ∞` suppresses any typ-
176 ing within the change-point window entirely. This parameter trades off two
177 failure modes that we quantify separately in Section 5: spurious type labels
178 from coordinated shifts in weakly correlated signals (reduced by stronger sup-
179 pression) against a masking dead-band for genuine faults co-occurring with a
180 change point (widened by stronger suppression). We treat `suppress_scale`
181 as the principal tuning knob of the composed system and report its effect
182 rather than fixing it silently.

183 4. Experimental Setup

184 We evaluate the composed system along two complementary axes: con-
185 trolled fault injection, which provides per-type ground truth, and a real fault
186 deployment, which provides an external validity check. Because no public
187 stream simultaneously labels all four fault types per signal and per timestep,
188 neither axis alone suffices, and we triangulate the two. All evaluation artifacts
189 referenced below are committed as CSV files under `examples/benchmarks/`.

190 4.1. Controlled fault injection

191 Following the injection-protocol precedent of Lai et al. (2021), we inject
192 parameterized faults of each type—bias, drift, accuracy loss, and freezing—
193 into a multivariate process with known ground truth, drawn from the Con-
194 trolled Anomalies Time Series (CATS) benchmark (Fleith, 2023). Each typed
195 fault is injected at controlled magnitude and onset onto a target signal,
196 and recall is measured per type. We run $N = 85$ trials per type in the
197 scaled controlled validation, and report Wilson 95% confidence intervals on
198 every recall and precision estimate (`i7_scaled_per_type.csv`). To probe
199 the dependence of bias recall on conditional coupling, we report bias re-
200 call per target channel against that channel’s conditional standard deviation
201 (`i7_scaled_bias_by_channel.csv`, $N = 17$ per channel). Freeze attribu-
202 tion versus base detection is evaluated on the same channels, recording per-
203 channel freeze and CDF detection rates and detection delays (`i7_freeze_latency.csv`).

204 4.2. Synthetic correlation control

205 To separate the two robustness mechanisms of the composed system,
206 we generate synthetic multivariate streams with controlled pairwise corre-
207 lation $\rho \in \{0.1, 0.3, 0.5, 0.7, 0.9\}$ and inject coordinated shifts of magnitude
208 $\{3, 6, 10\}$ conditional standard deviations co-occurring with a change point.
209 For each cell we sweep the suppression scale `suppress_scale` $\in \{1, 5, \infty\}$
210 (abbreviated `scale` hereafter) and report the mean false-label rate over 5
211 seeds (`i7_regime_fp.csv`). The co-occurring-fault dead-band is measured
212 by injecting a drift simultaneously with a change point and recording the
213 detection magnitude at each suppression scale (`i7_deadband.csv`). The
214 freeze-on-quiescent ablation injects no fault into a genuinely live but quiet sig-
215 nal (noise floor as a fraction of marginal standard deviation) and measures
216 the false freeze-label rate under default, strict, and loose freeze constants
217 (`i7_freeze_quiescent.csv`).

218 4.3. Real fault deployment

219 For external validity we use the Intel Berkeley Lab sensor deployment (Bodik
220 et al., 2004), in which motes report temperature and humidity over a diurnal
221 cycle. As a mote’s battery depletes (voltage below 2.3 V, held out as a physi-
222 cal fault label), its temperature and humidity readings degrade to physically
223 impossible values—temperatures of 100–122 °C and negative humidity. We
224 treat the battery-voltage label as ground truth for a real fault, never expos-
225 ing it to the detector or classifier. We sweep the detector’s mean-threshold
226 operating point and report, pooled over motes, the healthy false-positive
227 rate and the real-fault detection rate ($N_{\text{healthy}} \approx 1.52\text{M}$, $N_{\text{fault}} \approx 0.39\text{M}$;
228 `i7_intel_lab_operating_points.csv`), together with the distribution of
229 types assigned to detected real faults (`i7_intel_lab_fault_types.csv`).
230 To establish that online adaptation is required rather than optional, we ad-
231 ditionally fit a static model on a short prefix of the same stream and measure
232 its false-alarm rate.

233 4.4. Metrics

234 For the controlled injections we report per-type recall (fraction of injected
235 faults of a type that are detected) and cross-talk precision on normal peer
236 signals. For the real deployment we report the healthy false-positive rate and
237 the real-fault detection rate at each operating point. All proportions esti-
238 mated from finite samples are reported with Wilson 95% confidence intervals.

Table 1: Real battery-depletion fault detection on the Intel Berkeley Lab deployment, online adaptive detector, pooled over motes. Source: `i7_intel_lab_operating_points.csv` ($N_{\text{healthy}} = 1\,518\,828$, $N_{\text{fault}} = 387\,671$).

Mean threshold	Healthy FP rate	Real-fault detection
3.0	0.221	0.926
4.0	0.165	0.916
5.0	0.129	0.910
6.0	0.104	0.906
8.0	0.074	0.901

We do not report aggregate accuracy across types, because the type populations are not equiprobable in either the injection design or the real stream, and an aggregate would obscure the per-type and per-coupling structure that is the substance of our characterization.

5. Results

5.1. Real-fault operating curve

On the Intel Berkeley Lab deployment, the composed system detects real battery-depletion faults at a high rate while keeping healthy false alarms low, and the operator can trade the two by moving a single mean-threshold operating point. Table 1 reports the pooled sweep over motes ($N_{\text{healthy}} \approx 1.52\text{M}$ healthy samples, $N_{\text{fault}} \approx 0.39\text{M}$ fault samples; `i7_intel_lab_operating_points.csv`). At a mean threshold of 8.0 the system reaches a 7.4% healthy false-positive rate at 90.1% real-fault detection; lowering the threshold to 3.0 raises detection to 92.6% at the cost of a 22.1% false-positive rate. The detection rate is remarkably flat across the sweep (from 0.926 down to 0.901), so the operating point is chosen almost entirely by the tolerable false-alarm budget.

Online adaptation is mandatory on this stream, not a convenience. A static model fit on a short prefix of the same diurnal signal false-alarms at a rate of 0.97: the daily temperature and humidity cycle alone is enough to push a fixed model permanently out of distribution. The adaptive detector tracks the diurnal baseline and reserves its alarms for the battery-depletion excursions.

The types assigned to the detected real faults are dominated by freezing, as a stuck-at-degraded-value failure should be. Table 2 reports the type distribution over detected real-fault samples (`i7_intel_lab_fault_types.csv`):




Figure 1: Operating curve on the Intel Berkeley Lab battery-depletion faults (online adaptive detector, pooled over motes). Raising the mean threshold trades a lower healthy false-positive rate against a marginal drop in real-fault detection. Source: `i7_intel_lab_operating_points.csv`.

93.2% are typed freezing, 5.1% drift, 1.7% bias, and 0.06% accuracy loss. The freezing label corresponds to the physical reality of a transmitter pinned at a corrupted reading as its battery fails, so the typing layer converts the detector’s alarm into the maintenance-relevant diagnosis “replace the mote” rather than “investigate the process.”

The residual healthy false positives are themselves informative. Approximately 93% of healthy false positives are typed as drift, arising from the healthy diurnal ramps that the two-signal conditional model (temperature given humidity, and vice versa) cannot fully cancel. This is an honest limitation tied directly to the correlation-dependence we characterize in Section 5.4: with only two correlated signals available, slow shared ramps leak into the residual; a deployment with more correlated signals would suppress them further.

5.2. Per-type recall and the bias–coupling relationship

On the scaled controlled injections, the four types are recovered at sharply different rates. Table 3 reports per-type recall with Wilson 95% confidence

Table 2: Fault types assigned to detected real faults on the Intel Berkeley Lab deployment. Source: `i7_intel_lab_fault_types.csv`.

Assigned type	Count	Fraction
freezing	578 721	0.932
drift	31 727	0.051
bias	10 395	0.017
accuracy loss	343	0.001

Table 3: Per-type recall on scaled controlled injections, $N = 85$ trials per type, Wilson 95% CI. Source: `i7_scaled_per_type.csv`.

Fault type	Recall	95% CI	N
freezing	1.000	[0.957, 1.000]	85
drift	0.741	[0.639, 0.822]	85
accuracy loss	0.659	[0.553, 0.751]	85
bias	0.518	[0.413, 0.621]	85

intervals over $N = 85$ trials per type (`i7_scaled_per_type.csv`). Freezing is recovered perfectly (1.000, CI 0.957–1.000) and drift well (0.741), while accuracy loss (0.659) and especially bias (0.518) are lower. Cross-talk precision on normal peer signals is high: of the signals labelled by typing, the normal-peer label is correct with precision ≈ 0.91 (`i7_scaled_normal_precision.csv`), so coupled peers are rarely misattributed a fault.

The aggregate bias recall of 0.518 is not an unexplained weakness; it is the mean of two distinct regimes. Table 4 reports bias recall per target channel against that channel’s conditional standard deviation σ (`i7_scaled_bias_by_channel.csv`). Bias recall is monotone in σ : on strongly coupled channels (small conditional σ , because peers explain most of the variance) the recall is near-perfect (0.941 at $\sigma = 0.216$, 1.000 at $\sigma = 0.480$), whereas on weakly coupled channels (large conditional σ) it collapses (0.353, 0.235, down to 0.059 at $\sigma = 1.471$). This is the key finding of the characterization: the interpretable conditional model *absorbs* a constant offset on a strongly coupled signal, because the conditional mean tracks the offset through the peers, while it *isolates* the same offset on a weakly coupled signal. The aggregate 0.52 is thus the average of an “isolable” regime and an “absorbed” regime, not a uniformly mediocre detector.

Table 4: Bias recall is monotone in conditional standard deviation σ : offsets are absorbed on strongly coupled (small- σ) channels and isolated on weakly coupled (large- σ) channels. $N = 17$ per channel, Wilson 95% CI. Source: `i7_scaled_bias_by_channel.csv`.

Channel	Conditional σ	Bias recall	95% CI	N
bed2	0.216	0.941	[0.73, 0.99]	17
bed1	0.480	1.000	[0.816, 1.00]	17
bso1	1.439	0.353	[0.173, 0.587]	17
bfo2	1.443	0.235	[0.096, 0.473]	17
cso1	1.471	0.059	[0.01, 0.27]	17

5.3. Freezing is type attribution, not detection capability

We initially expected freezing to be a detection blind spot of the base detector, on the reasoning that a signal stuck at a value within its normal range would not trigger a conditional-CDF alarm. This claim is refuted by our own validation. Table 5 reports, per channel, the base CDF detection rate, the dedicated freeze-test detection rate, their median detection delays, and the fraction of trials in which the freeze test fired earlier (`i7_freeze_latency.csv`). The base CDF detector catches *every* injected freeze on *every* channel (`cdf_detect_rate` = 1.0, `frac_cdf_never` = 0.00 throughout), because the conditional expectation keeps moving while the frozen signal does not, driving the residual out of bounds. Moreover, the CDF detector is usually *faster*: its median delay ranges from 2 to 28 steps, and the freeze test—whose fixed 25-step confirmation window sets its median delay at exactly 25—fires earlier in only about half the trials, and only on the most isolated channel (bed2, `frac_freeze_earlier` = 0.529); on the most coupled channels it never beats the CDF detector (0.0).

The conclusion is that the freeze test adds type *attribution*—distinguishing a stuck sensor from a genuine process excursion—rather than detection capability or speed. The theoretical blind spot, in which the conditional standard deviation collapses to zero and the residual cannot grow, is an edge case that does not arise in the real or injected data we examined. On the real deployment this attribution is exactly what types the battery-depletion faults as freezing (Section 5.1), turning a detection into a maintenance label.

5.4. Regime robustness comes from suppression, not residual cancellation

We also expected the conditional residual to cancel coordinated shifts—a fault arriving as a shared shift across correlated signals would, on this




Figure 2: Bias recall is monotone in conditional σ : a constant offset is isolated on weakly coupled (small- σ) channels and absorbed by the conditional mean on strongly coupled (large- σ) channels. Error bars are Wilson 95% intervals. Source: `i7_scaled_bias_by_channel.csv`.

reasoning, leave the conditional residual unchanged and so produce no spurious type label, independent of the change-point suppression. This claim is also refuted: residual cancellation (mechanism 1) is strongly correlation-dependent, and the robustness of the composed system instead comes from graded change-point suppression (mechanism 2). Table 6 reports the false-label rate across correlation ρ , coordinated-shift magnitude, and suppression scale (`i7_regime_fp.csv`, 5 seeds per cell). With suppression disabled (`scale = 1`), a $\geq 6\sigma$ coordinated shift at low correlation ($\rho = 0.1$) produces about 20% false labels (0.199 at 6σ , 0.201 at 10σ); the boundary recedes as ρ increases and is gone by $\rho \geq 0.7$. By contrast, the default graded suppression (`scale = 5`) drives the false-label rate to exactly 0.0 in every (ρ, shift) cell, as does full suppression (`scale = \infty`).

The adapt-into-drift robustness of the composed system therefore rests on mechanism 2, the graded change-point suppression, and not on the conditional-residual cancellation we originally credited. This matters operationally: the suppression scale is the parameter the operator must set, and it has a mea-

Table 5: Freezing detection: the base CDF detector catches every injected freeze on every channel, usually faster than the fixed 25-step freeze test. Source: `i7_freeze_latency.csv` ($N = 17$ per channel).

Channel	Cond. σ	CDF detect	CDF median delay	Freeze earlier	CDF never
bed2	0.216	1.0	28	0.529	0.0
bed1	0.480	1.0	9	0.294	0.0
bso1	1.439	1.0	17	0.412	0.0
bfo2	1.443	1.0	3	0.000	0.0
cso1	1.471	1.0	2	0.000	0.0

341 surable cost, quantified next.

342 5.5. The cost of suppression: dead-band and freeze-on-quiescent ablations

343 Suppression suppresses real faults too. Table 7 reports the detection of a
 344 genuine drift that co-occurs with a change point, as a function of suppression
 345 scale (`i7_deadband.csv`, 5 seeds). With no suppression (`scale = 1`) the
 346 drift is detected once it reaches 5.2σ ; at the default `scale = 5` detection is
 347 delayed until 17.6σ ; at `scale =  $\infty$` the drift is never detected. The default
 348 suppression thus opens a masking dead-band of roughly $[3\sigma, 15\sigma]$ for faults
 349 that happen to coincide with a change point—the measured price of the
 350 robustness reported in Section 5.4.

351 The freeze test similarly requires calibration on steady-state signals. Ta-
 352 ble 8 reports the false freeze-label rate on a genuinely live but very quiet
 353 signal (noise at 0.5% of the marginal standard deviation) under three freeze-
 354 constant settings (`i7_freeze_quiescent.csv`, 5 seeds). With the default
 355 constants, such a signal is falsely labelled frozen 72% of the time; the loose
 356 constants make it worse ($\approx 88\%$); the strict constants (`freeze_eps =  $10^{-4}$` ,
 357 `freeze_var_ratio =  $5 \times 10^{-3}$` , `freeze_abs_scale = 5`) eliminate the false
 358 label entirely (0.0). We therefore recommend the strict constants, or per-
 359 signal calibration, for deployments containing steady-state signals.

360 6. Discussion

361 The characterization in Section 5 maps the operating envelope of the
 362 composed system honestly. We read each apparent limitation as a quantified
 363 property that an operator can plan around, rather than a defect to be hidden.

Table 6: False-label rate under coordinated shifts co-occurring with a change point, by correlation ρ , shift magnitude, and suppression scale (mean over 5 seeds). Residual cancellation alone (`scale = 1`) fails at low correlation; graded suppression (`scale = 5`) yields 0.0 everywhere. Source: `i7_regime_fp.csv`.

ρ	Shift (σ)	<code>scale=1</code>	<code>scale=5</code>	<code>scale=∞</code>
0.1	3	0.000	0.000	0.000
0.1	6	0.199	0.000	0.000
0.1	10	0.201	0.000	0.000
0.3	3	0.000	0.000	0.000
0.3	6	0.021	0.000	0.000
0.3	10	0.198	0.000	0.000
0.5	3	0.000	0.000	0.000
0.5	6	0.000	0.000	0.000
0.5	10	0.112	0.000	0.000
0.7	3–10	0.000	0.000	0.000
0.9	3–10	0.000	0.000	0.000

Table 7: Co-occurring-fault dead-band: detection magnitude of a drift co-occurring with a change point, by suppression scale (5 seeds). Source: `i7_deadband.csv`.

Suppression scale	Detected	Median delay	Magnitude at detection (σ)
1.0	yes	26	5.2
5.0	yes	88	17.6
∞	no	—	—

364 *Bias absorption is a property of the interpretable model, not a bug..* The
365 monotone dependence of bias recall on conditional standard deviation (Ta-
366 ble 4) follows directly from what makes the base detector interpretable. Be-
367 cause the conditional mean of a strongly coupled signal tracks its peers, a
368 constant offset injected on that signal is partly explained away by the condi-
369 tional model, and the residual—hence the typing—does not see it. The same
370 mechanism that yields tight, peer-informed operating limits also absorbs off-
371 sets on those signals. The practical reading is that bias is reliably typed
372 where it is observable (weakly coupled signals show recall up to 0.941) and is
373 absorbed where the model is confident enough to overrule it; the aggregate
374 0.518 is the average of these two regimes, not a uniform failure. An oper-




Figure 3: Without change-point suppression, conditional-residual cancellation (mechanism 1) fails at low correlation: a large coordinated shift produces false fault labels until ρ is high. The default graded suppression (mechanism 2) drives the rate to zero in every cell. Source: `i7_regime_fp.csv`.

375 ator who needs bias coverage on a strongly coupled signal should monitor
 376 that signal’s raw calibration directly, a check the conditional model is, by
 377 construction, not designed to provide.

378 *Two-signal cancellation limits the real-deployment false-positive floor..* On
 379 the Intel Lab stream the residual healthy false positives are $\approx 93\%$ drift,
 380 traced to diurnal ramps that a two-signal conditional model cannot fully
 381 cancel. This is the same correlation-dependence quantified in Table 6: with
 382 few correlated signals, slow shared movement leaks into the residual. The
 383 limitation is therefore structural and predictable—it improves with the num-
 384 ber of correlated signals available—rather than a tuning failure. We report
 385 the 7.4% floor at 90% detection as the achievable operating point on a two-
 386 signal deployment, with the explicit expectation that richer instrumentation
 387 would lower it.

388 *The dead-band is the priced cost of adaptation robustness..* Because robust-
 389 ness to faults arriving during adaptation comes from graded change-point




Figure 4: The co-occurring-fault dead-band: a drift co-occurring with a change point must grow from 5.2σ (no suppression) to 17.6σ (default suppression) before detection, and is never detected under full suppression. Source: `i7_deadband.csv`.

Table 8: Freeze-on-quiescent ablation: false freeze-label rate on a live but quiet signal (noise 0.5% of marginal std), by freeze-constant setting (5 seeds). Source: `i7_freeze_quiescent.csv`.

Freeze constants	<code>freeze_eps</code>	<code>freeze_var_ratio</code> / <code>freeze_abs_scale</code>	False freeze rate
default	1×10^{-3}	1×10^{-2} / 20	0.717
loose	5×10^{-3}	5×10^{-2} / 50	0.887
strict	1×10^{-4}	5×10^{-3} / 5	0.000

390 suppression and not from residual cancellation (Section 5.4), it carries a
391 cost: the masking dead-band of roughly $[3\sigma, 15\sigma]$ for faults co-occurring with
392 a change point (Table 7). We present this as an explicit trade-off governed
393 by a single parameter, `suppress_scale`. A deployment that fears spurious
394 labels from coordinated shifts should keep the default suppression; one that
395 fears missing a fault during a regime change should lower it, accepting the
396 false-label rate of Table 6. The point is that the trade-off is measured and
397 tunable, not silent.

398 *Threshold and freeze-constant calibration are deployment decisions..* The operating-
399 point sweep (Table 1) and the freeze-on-quiescent ablation (Table 8) both
400 show that the system’s behavior is set by a small number of interpretable con-

401 stants. The mean threshold selects the false-alarm/detection trade-off; the
 402 freeze constants determine how aggressively a quiet signal is judged frozen,
 403 with the strict setting eliminating false freeze labels on a live-but-quiet signal
 404 that the default mislabels 72% of the time. We recommend the strict freeze
 405 constants or per-signal calibration for steady-state signals, and we expose the
 406 threshold as the operator’s primary control. These are calibration decisions
 407 with quantified consequences, not hidden defaults.

408 *Data availability constrains the evaluation to triangulation..* No public stream
 409 simultaneously labels all four fault types per signal and per timestep, which
 410 is why we combine controlled injection (per-type ground truth, but syn-
 411 thetic) with a real fault deployment (genuine faults, but a single physical
 412 failure mode with a held-out label). Neither alone would be conclusive: the
 413 injections could overstate performance on idealized faults, and the real de-
 414 ployment exercises mainly the freezing signature. Read together, the injec-
 415 tions establish the per-type envelope and the refuted claims, while the real
 416 deployment confirms that the envelope holds on genuine faults at a usable
 417 operating point. We regard this triangulation as the appropriate evidentiary
 418 standard given the data that exist, and we flag the absence of a fully labelled
 419 multi-type stream as the principal external constraint on stronger claims.

420 7. Conclusion

421 We presented a streaming sensor-fault classifier that closes the detection-
 422 to-action gap of a self-supervised online anomaly detector (Wadinger and
 423 Kvasnica, 2023) by attributing each flagged signal to a maintenance-relevant
 424 fault type—bias, drift, accuracy loss, or freezing—from the detector’s own
 425 conditional residuals, in $O(1)$ memory and with no history buffers. The
 426 fault taxonomy is the established one (Sharma et al., 2010; Ni et al., 2009);
 427 the contribution is its streaming, conditional-residual realization composed
 428 directly with the detector, so that type attribution is available at the same
 429 latency and operational footprint as detection.

430 The headline result is practical. On real battery-depletion faults in the
 431 Intel Berkeley Lab deployment (Bodik et al., 2004), the composed system
 432 reaches 90% real-fault detection at a 7.4% healthy false-positive rate and
 433 types the detected faults predominantly as freezing (0.93), the correct main-
 434 tenance diagnosis for a transmitter stuck at a degraded value. Online adap-
 435 tation is shown to be mandatory: a static prefix model false-alarms at 0.97
 436 on the same diurnal stream.

437 Around this result we characterized the operating envelope rather than
438 overstating it. Freezing and drift are recovered with recall 1.000 and 0.741 on
439 controlled injections, while bias recall is monotone in conditional coupling—
440 absorbed on strongly coupled signals and isolated on weakly coupled ones—so
441 that the aggregate 0.518 is the mean of two explained regimes. We reported
442 two refuted claims as such: freezing is not a detection blind spot of the
443 base detector (which catches every injected freeze, usually faster than the
444 dedicated test, so the freeze test contributes type attribution rather than
445 detection), and adaptation robustness arises from graded change-point sup-
446 pression rather than from conditional-residual cancellation. The suppression
447 carries a measured masking dead-band of roughly $[3\sigma, 15\sigma]$ for co-occurring
448 faults, and the freeze test requires strict constants or per-signal calibration
449 on steady-state signals.

450 The scope is correspondingly honest. The real-deployment evidence exer-
451 cises principally the freezing signature, the two-signal conditional model sets
452 a structural false-positive floor that richer instrumentation would lower, and
453 per-type ground truth is available only through injection because no public
454 stream labels all four fault types per signal and per timestep. Within that
455 scope, the system delivers a deployable typing layer that turns an adaptive
456 detector’s alarm into an actionable maintenance label, with an operating
457 envelope that is mapped rather than assumed.

458 **Additional information**

459 The classifier and all evaluation artifacts are openly accessible on GitHub
460 at the following URL: [https://github.com/MarekWadinger/online_outlier_](https://github.com/MarekWadinger/online_outlier_detection)
461 [detection](https://github.com/MarekWadinger/online_outlier_detection).

462 **CRediT authorship contribution statement**

463 **Marek Wadinger:** Conceptualization; Data curation; Formal analysis;
464 Investigation; Methodology; Software; Validation; Visualization; Writing -
465 original draft; and Writing - review & editing. **Michal Kvasnica:** Concep-
466 tualization; Funding acquisition; Project administration; Resources; Super-
467 vision; Validation.

468 Declaration of Competing Interest

469 The authors declare that they have no known competing financial inter-
470 ests or personal relationships that could have appeared to influence the work
471 reported in this paper.

472 Acknowledgements

473 The authors gratefully acknowledge the contribution of the Slovak Re-
474 search and Development Agency under the project APVV-20-0261. The au-
475 thors gratefully acknowledge the contribution of the Scientific Grant Agency
476 of the Slovak Republic under the grant 1/0490/23. This paper was funded
477 by the European Union under Horizon Europe Grant Agreement number
478 101079342 (Fostering Opportunities Towards Slovak Excellence in Advanced
479 Control for Smart Industries).

480 References

- 481 Bodik, P., Hong, W., Guestrin, C., Madden, S., Paskin, M., Thibaux, R.,
482 2004. Intel berkeley research lab sensor data. <https://db.csail.mit.edu/labdata/labdata.html>.
483
- 484 Fleith, P., 2023. Controlled anomalies time series (cats) dataset. doi:10.
485 5281/zenodo.8338435. solenix Engineering GmbH; CC-BY-4.0.
- 486 Lai, K.H., Zha, D., Xu, J., Zhao, Y., Wang, G., Hu, X., 2021. Revisiting time
487 series outlier detection: Definitions and benchmarks, in: Advances in Neu-
488 ral Information Processing Systems (NeurIPS) Datasets and Benchmarks
489 Track. URL: <https://openreview.net/forum?id=r8Iv0snHchr>.
- 490 Ni, K., Ramanathan, N., Chegade, M.N.H., Balzano, L., Nair, S., Zahedi, S.,
491 Kohler, E., Pottier, G., Hansen, M., Srivastava, M., 2009. Sensor network
492 data fault types. ACM Transactions on Sensor Networks 5, 1–29. doi:10.
493 1145/1525856.1525863.
- 494 Sharma, A.B., Golubchik, L., Govindan, R., 2010. Sensor faults: Detection
495 methods and prevalence in real-world datasets. ACM Transactions on
496 Sensor Networks 6, 1–39. doi:10.1145/1754414.1754419.

497 Wadinger, M., Kvasnica, M., 2023. Real-time outlier detection with dynamic
498 process limits, in: 2023 24th International Conference on Process Control
499 (PC), pp. 138–143. doi:10.1109/PC58330.2023.10217717.